

# Deep dive into Linux Hibernation: From kernel flows to Modern tools

Presented by

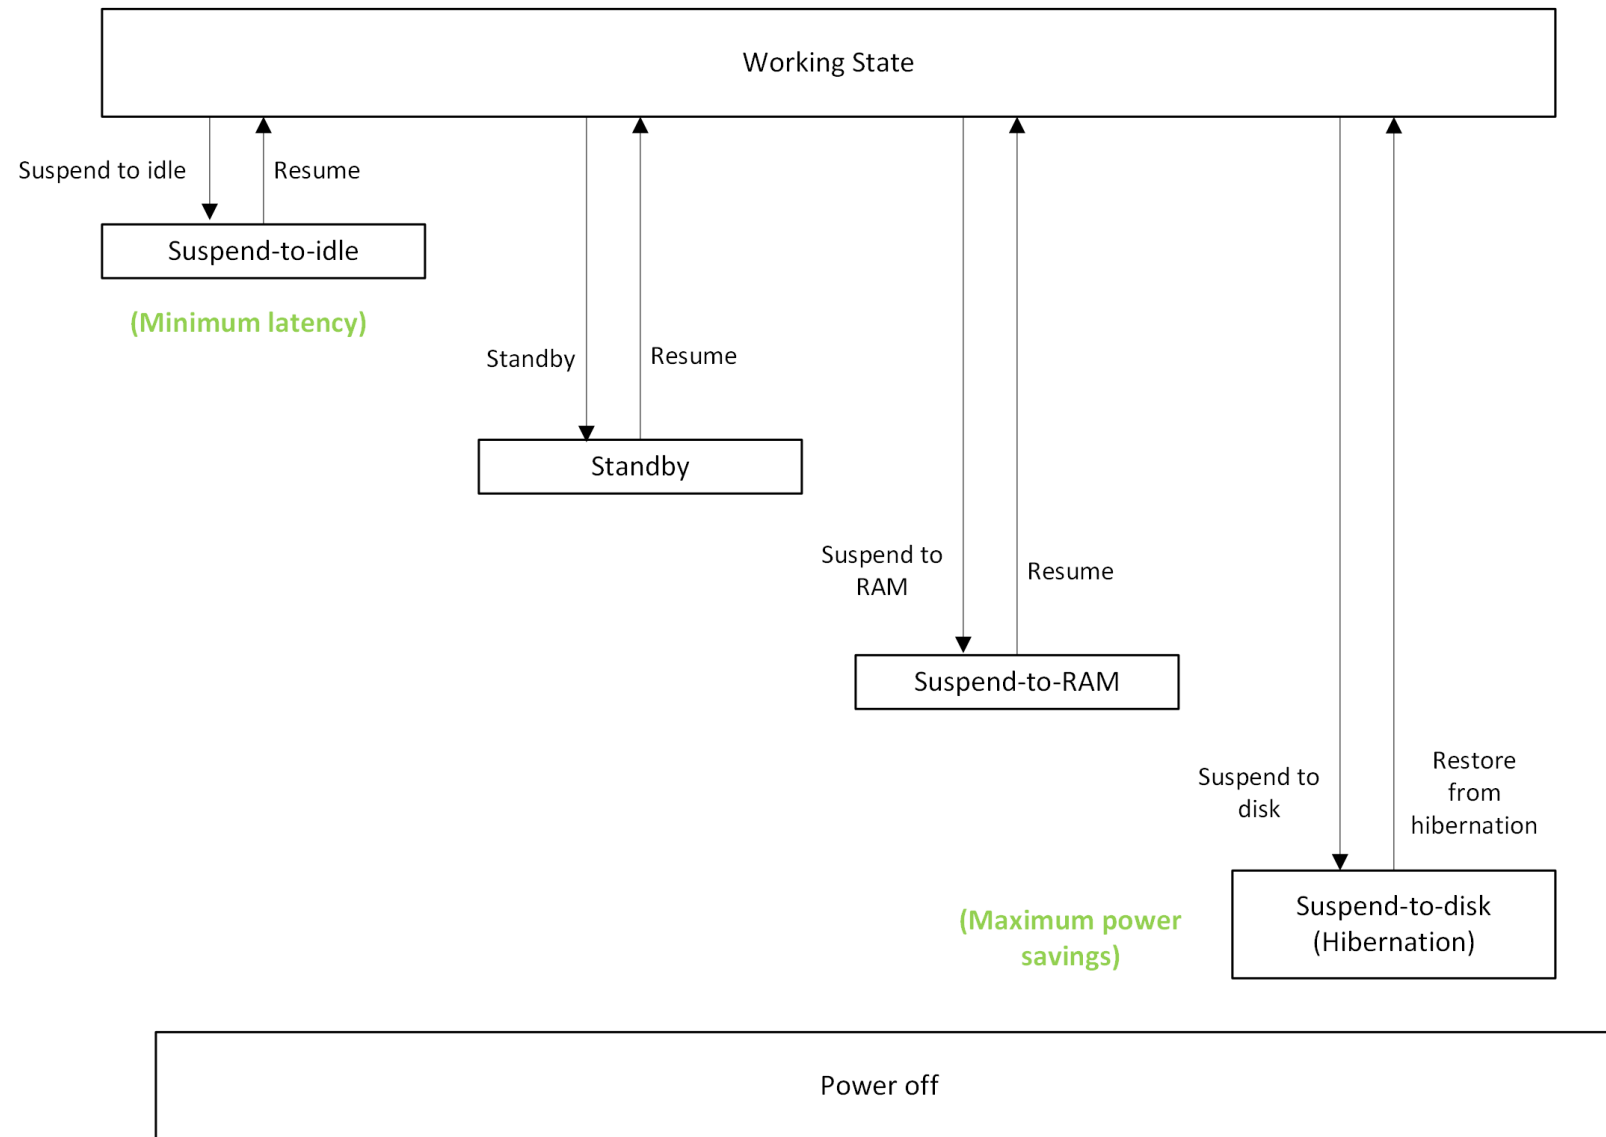
**Ankita Pareek** (Software Engineer, Microsoft)

**Erni Sri Satya Vennela** (Software Engineer, Microsoft)

# Contents

- Power Management in the Linux kernel
- Why hibernation?
- Understanding Linux Hibernation
- Deep dive into the hibernation implementation
- Sysfs interface
- Hibernation setup tool
- Critical Challenges in Hibernation Systems
- References

# Power Management States in the Linux kernel



# System Power Management Sleep States


Features	Suspend-to-Idle (S0ix)	Standby (S1)	Suspend-to-RAM (S3)	Hibernation (S4)
<b>What Happens</b>	CPU enters deepest idle state while keeping system in working state	CPU and some devices powered down while memory remains on	Most components powered off except memory which preserves system state	System state saved to disk and system powered off completely
<b>ACPI State</b>	S0ix	S1	S3	S4
<b>Power Consumption</b>	Highest	<S0ix but >S3	<S1 but >S4	Lowest
<b>Wake-up Time</b>	Minimum	>S0ix but <S3	>S1 but <S4	Maximum
<b>CPU State</b>	<ul style="list-style-type: none"> <li>- Deepest idle state</li> <li>- Core voltage reduced</li> <li>- Most clocks stopped</li> <li>- Caches maintained</li> </ul>	<ul style="list-style-type: none"> <li>- Non-boot CPU stopped</li> <li>- Context maintained</li> <li>- Some clocks running</li> <li>- Voltage still present</li> </ul>	<ul style="list-style-type: none"> <li>- Powered off</li> <li>- Context saved to RAM</li> <li>- All clocks stopped</li> <li>- Voltage off</li> </ul>	<ul style="list-style-type: none"> <li>- Completely powered off</li> <li>- Context saved to disk</li> <li>- Mostly all power rails off</li> </ul>
<b>Memory State</b>	<ul style="list-style-type: none"> <li>- Fully operational</li> <li>- Running at low power</li> <li>- All contents preserved</li> <li>- Memory controller active</li> </ul>	<ul style="list-style-type: none"> <li>- Fully powered</li> <li>- Normal operation</li> <li>- Full refresh rate</li> <li>- All contents preserved</li> </ul>	<ul style="list-style-type: none"> <li>- Self-refresh mode</li> <li>- Minimal power mode</li> <li>- Contents preserved</li> </ul>	<ul style="list-style-type: none"> <li>- Completely off</li> <li>- Contents saved to disk</li> <li>- No power consumption</li> </ul>
<b>Device State</b>	<ul style="list-style-type: none"> <li>- Devices in low power</li> <li>- Some peripherals suspended</li> <li>- Wake-capable devices active</li> </ul>	<ul style="list-style-type: none"> <li>- Some devices active</li> <li>- Most in low power mode</li> <li>- Clocks may be running</li> </ul>	<ul style="list-style-type: none"> <li>- Wake devices in low power</li> <li>- Most devices are suspended</li> <li>- Minimal power state</li> </ul>	<ul style="list-style-type: none"> <li>- All devices off</li> <li>- State saved to disk</li> <li>- Minimum power consumption</li> </ul>
<b>Debug Interface</b>	/sys/power/mem_sleep /sys/power/pm_debug	/sys/power/state /proc/acpi/sleep	/sys/power/state /sys/power/mem_sleep	/sys/power/disk /sys/power/state

Source: [System Sleep States — The Linux Kernel documentation](#)

# Why hibernation?



**Minimum Power  
Consumption**



**Cost Efficient**



**State Preservation**

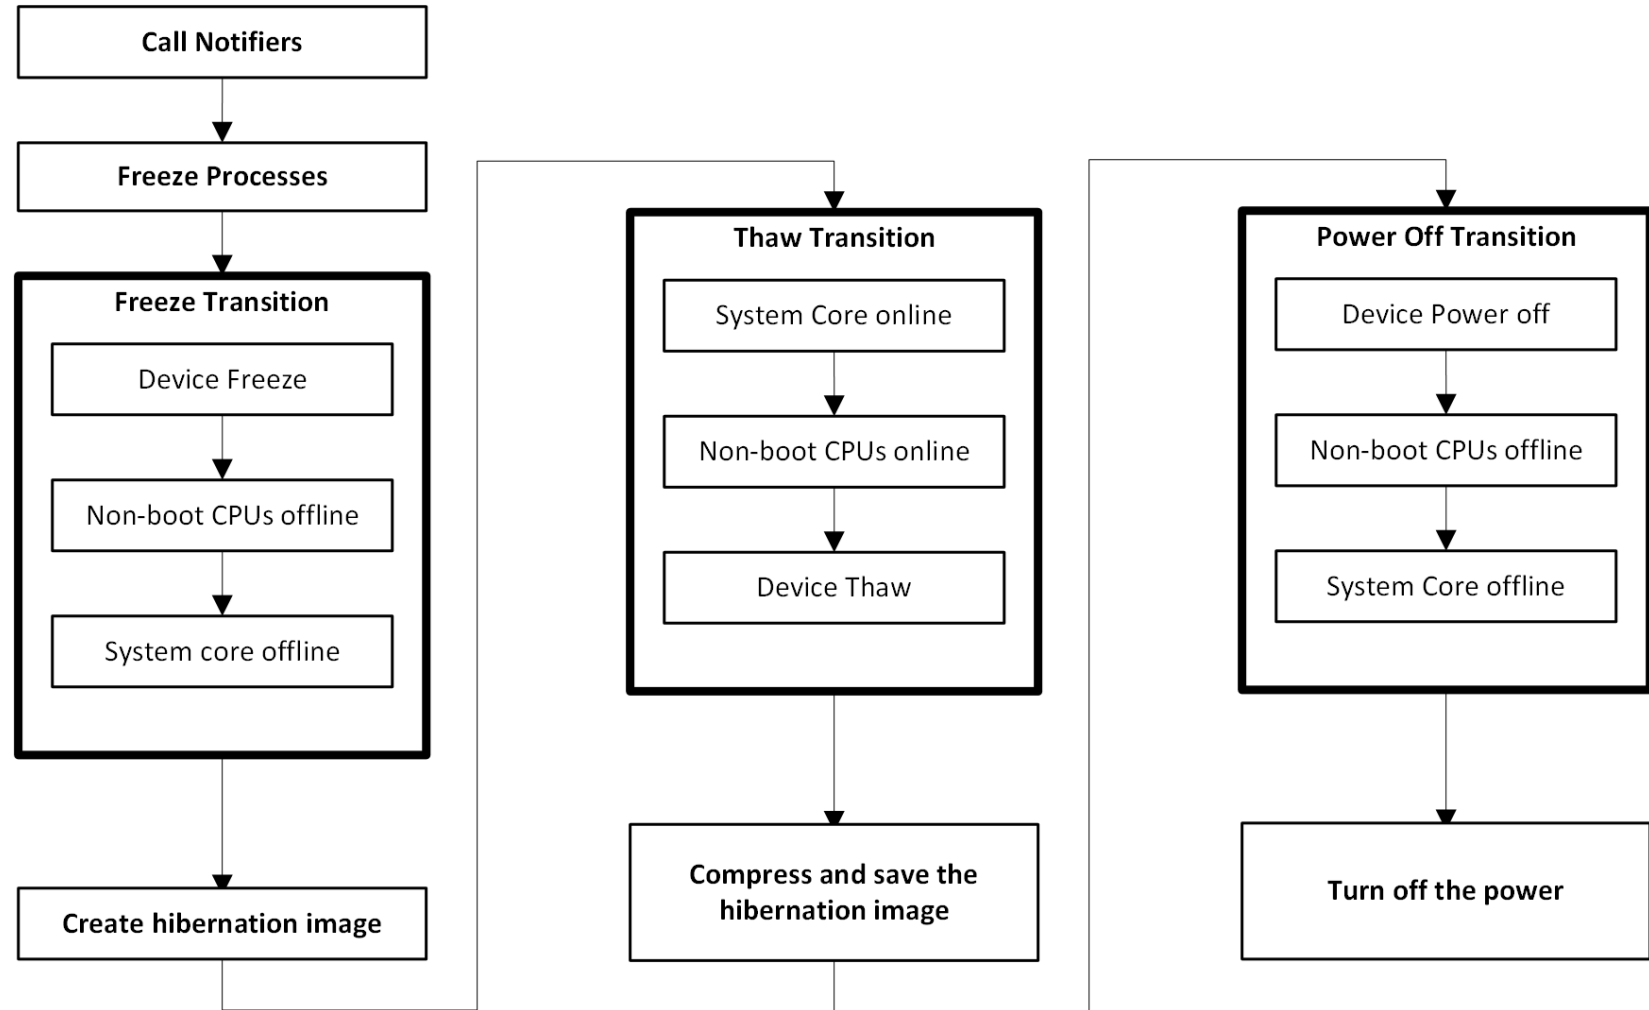


**Faster boot time**

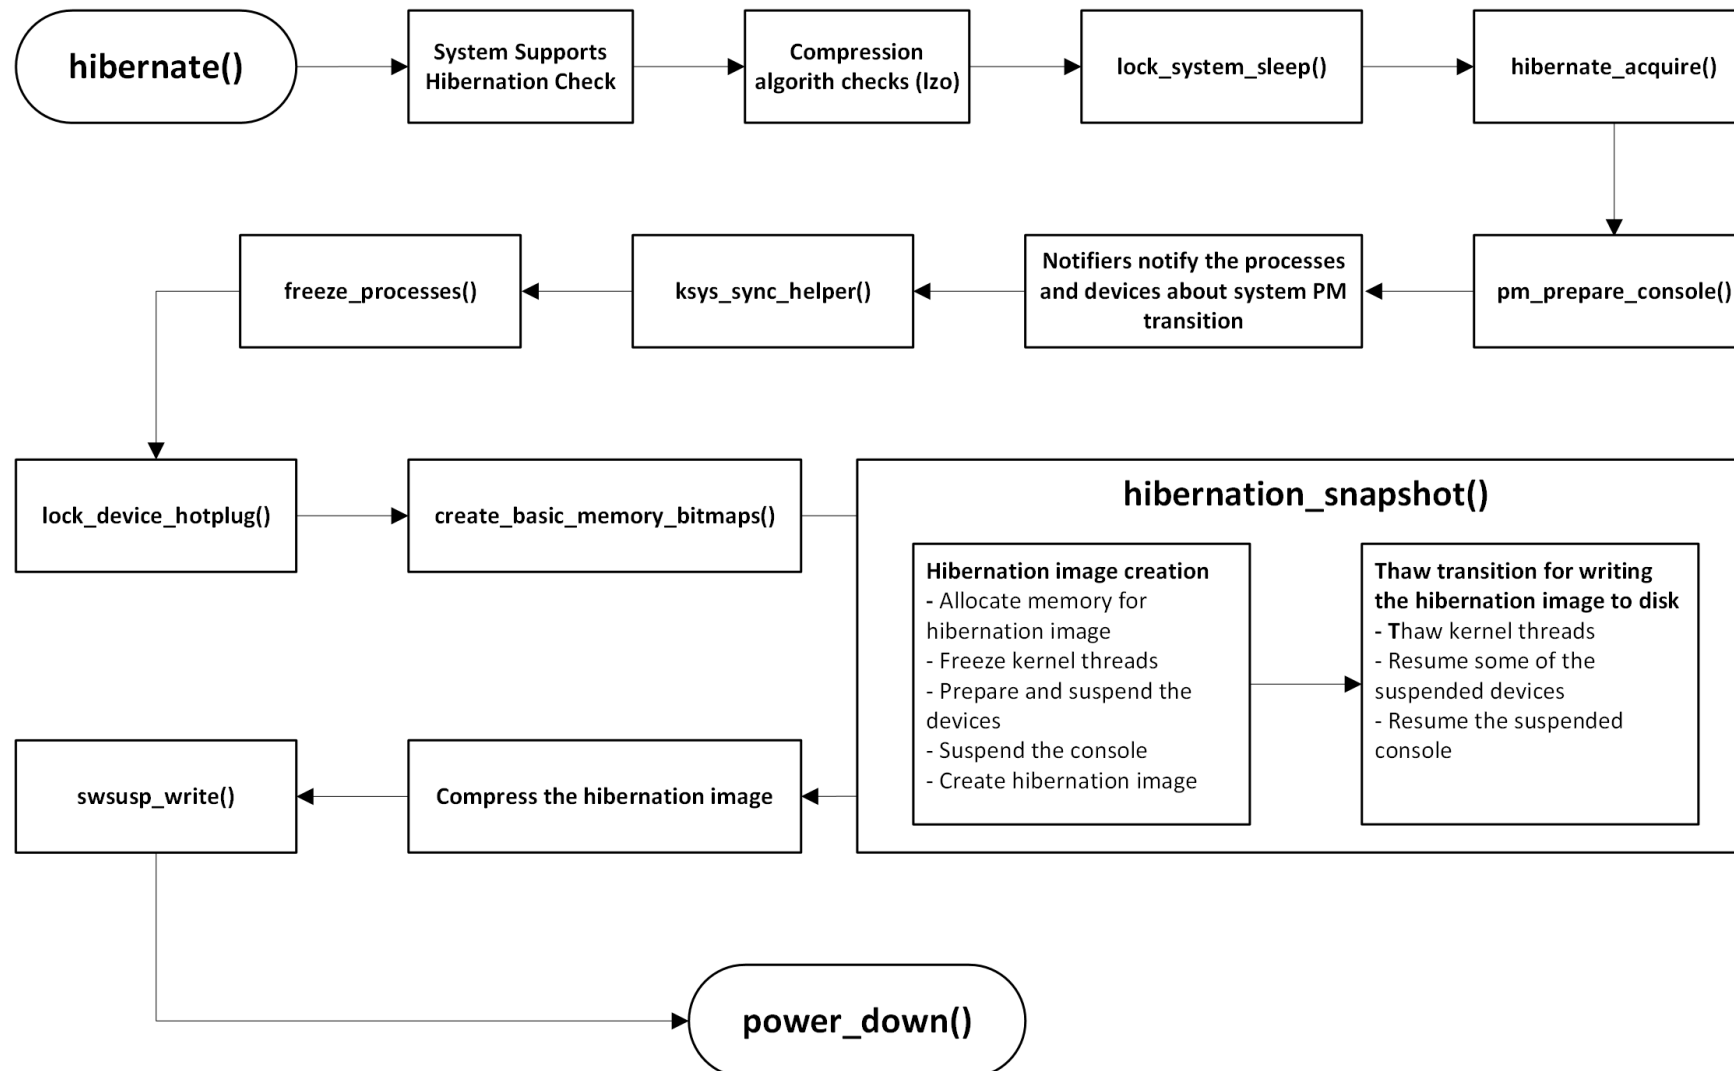
# Understanding Linux hibernation

- A system power-saving state aka Suspend-to-disk
- Saves complete system state to a non-volatile memory (disk/swap)
- System enters a special low-power state
- Restores system state on next boot
- Kernel config: **CONFIG\_HIBERNATION**
- Disabled when secure boot or kernel\_lockdown feature are enabled on the system
- Commonly supported on almost all Linux distros

# Code walkthrough (Basic Flow)

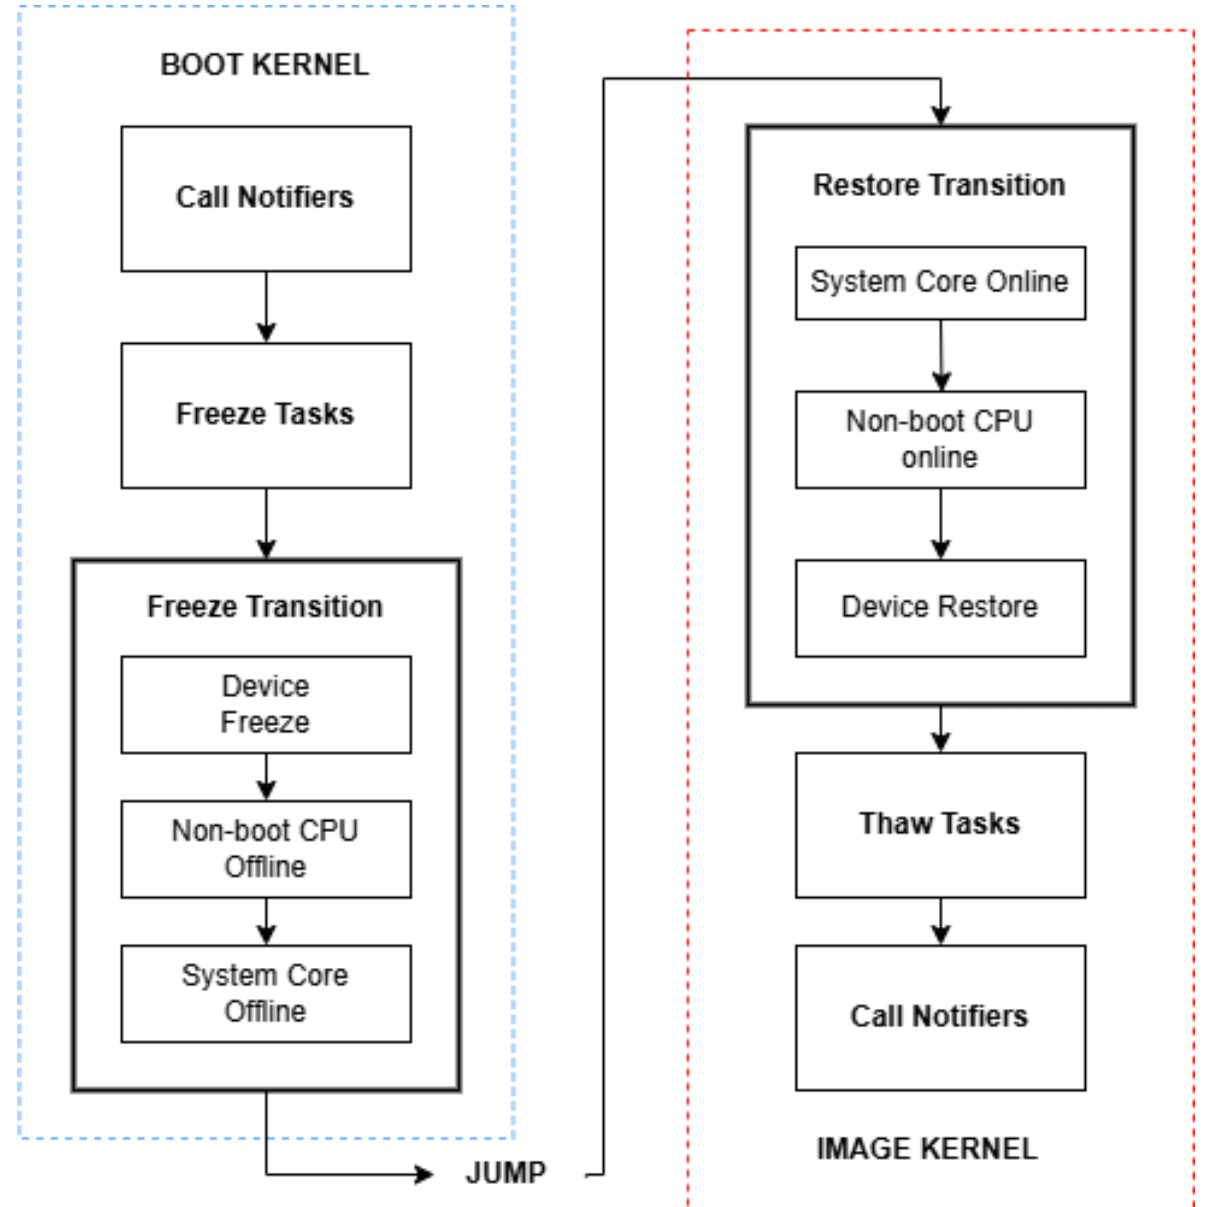


# Code walkthrough (Detailed Flow)

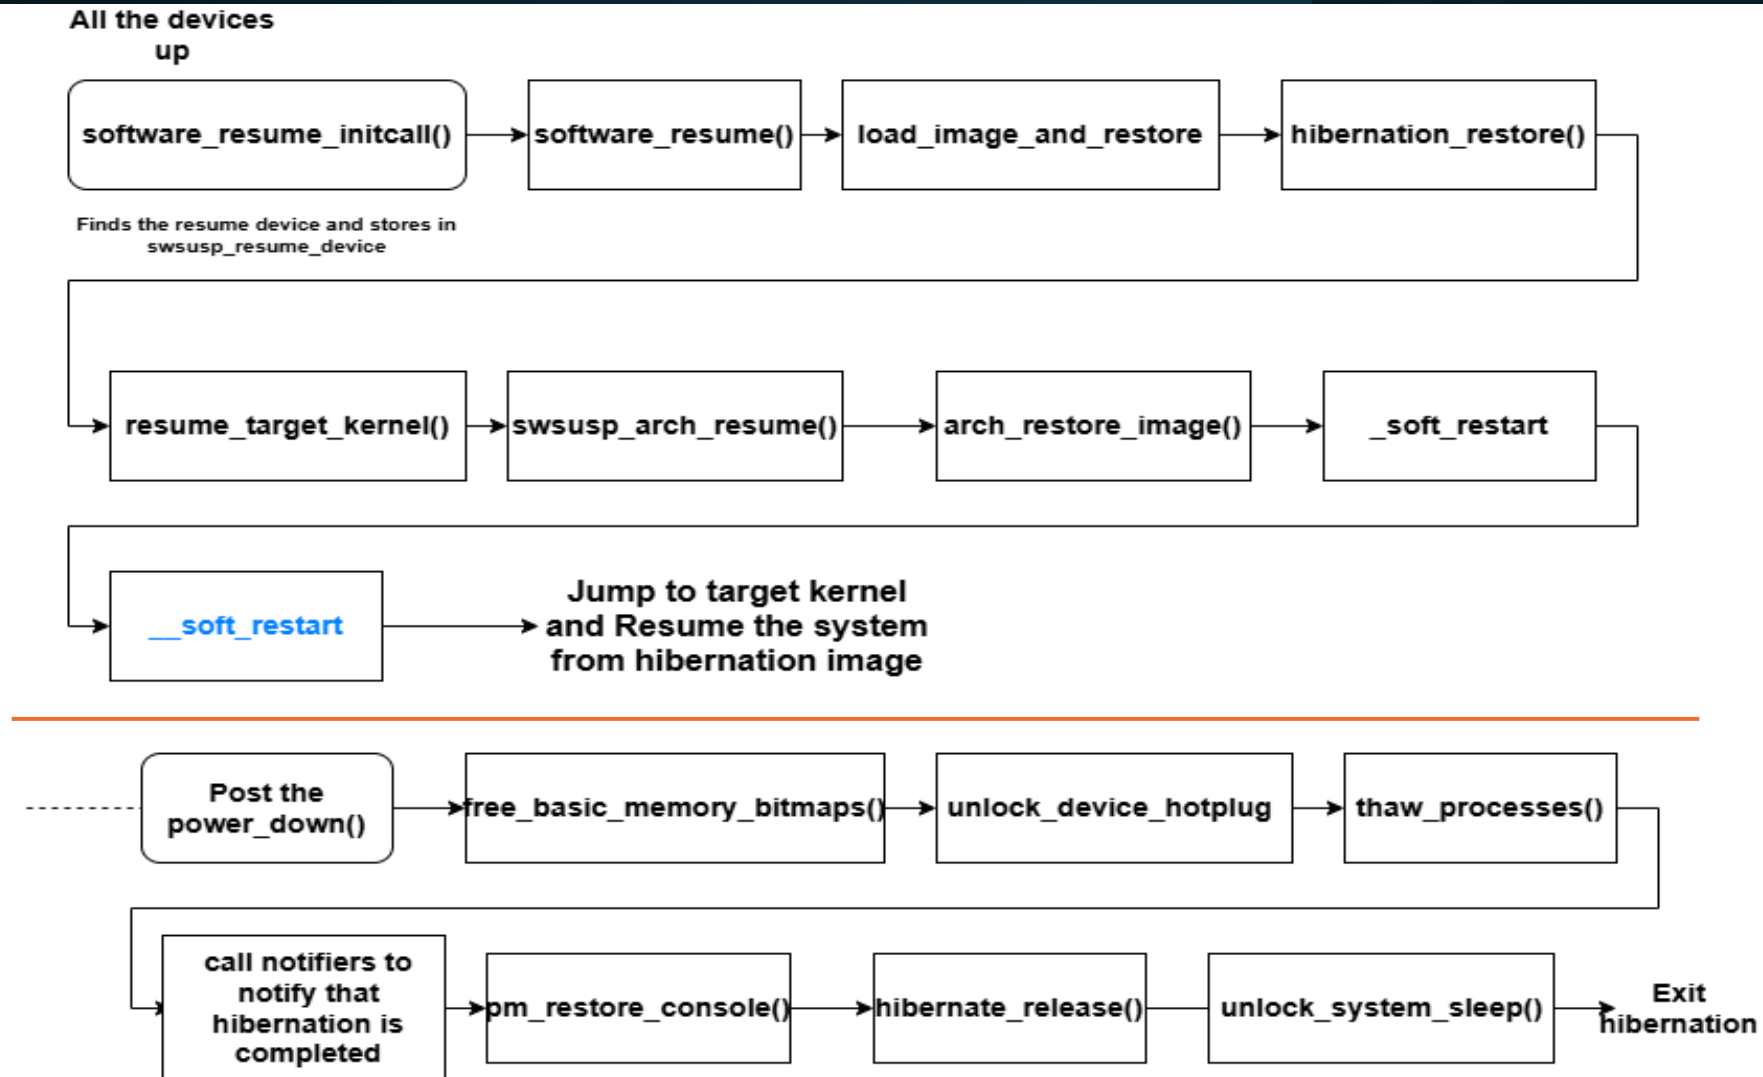




# Code Walkthrough (Resume Flow)



# Resume flow (detailed)



# Sysfs interface

The interface exists in /sys/power directory

**/sys/power/state:**

This controls the system power state

Reading from this file returns what sleep states are supported:

- **standby:** Power-on Suspend
- **freeze:** Suspend-to-idle
- **mem:** Suspend-to-RAM
- **disk:** Suspend-to-disk (hibernate)

Writing one of these strings to this file causes the system to transition into that state

# Sysfs interface

## **/sys/power/disk:**

Controls the operating mode of the Hibernation mechanism.

Reading from this file returns supported operating modes and the currently selected ones in brackets:

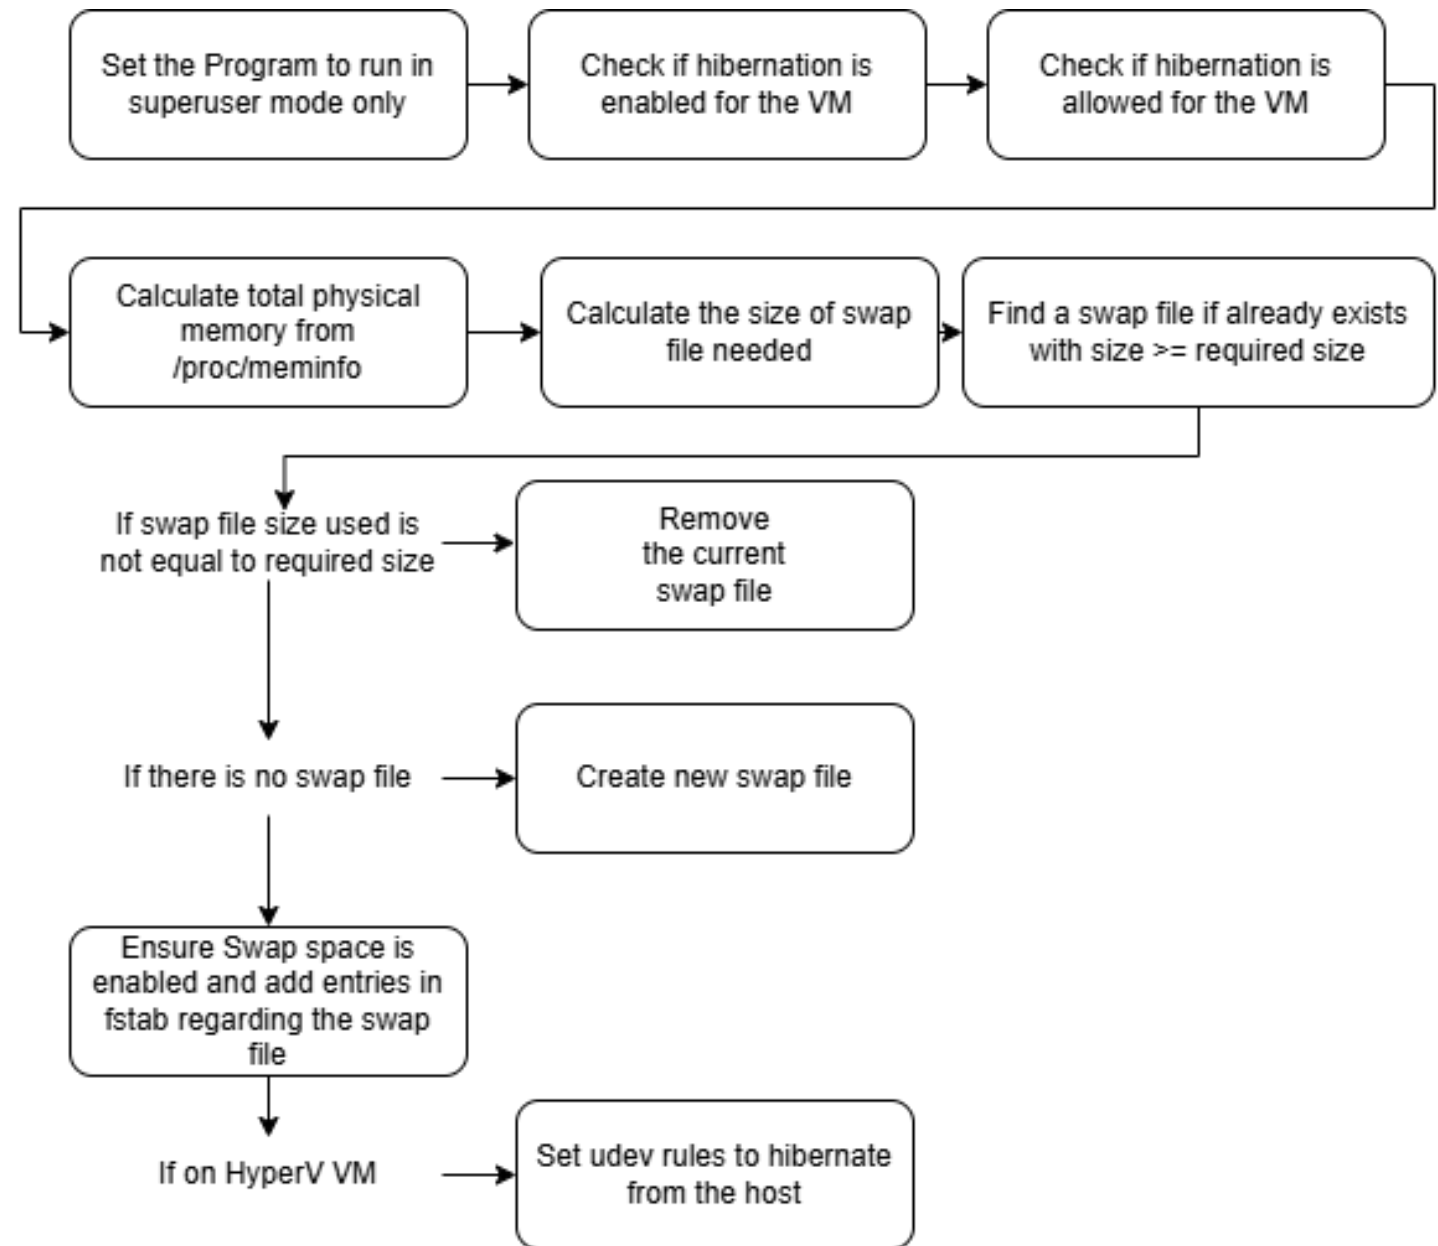
- **platform:** Default mode - uses platform driver for hibernation
- **shutdown:** Powers off the system after hibernation image creation without platform-specific calls
- **reboot:** Reboot the system after hibernation image creation
- **suspend:** Creates hibernation image but doesn't power down - keeps system in low-power state
- **test\_resume:** Test mode - creates and resumes from hibernation image without powering down
- **noresume:** Prevents resume from hibernation image on next boot

Writing one of these strings to this file will select the corresponding mode for hibernation

# Hibernation setup tool

- An open-source tool developed by Microsoft to configure and enable hibernation in Linux
- Creates optimized swap file according to the system RAM
- Configures kernel and boot parameters for immediate hibernation and resume operations
- Primarily supports distributions using GRUB2 bootloader and initramfs-tools
- Uses system hooks to track hibernation success, failures, and cold boot scenarios and store them in the system logs
- This is specifically designed for hibernating Linux VMs.

# How does this tool work?





Let us hibernate a VM using the hibernation-setup-tool!

# Challenges with hibernation



**MAX WAKE-UP  
TIMES**



**STORAGE  
REQUIREMENT**



**MEMORY STATE  
INCONSISTENCY**



**PLATFORM  
DEPENDENCIES**



**ENCRYPTED  
SYSTEMS**



# References

- [Power Management In The Linux\\* Kernel - Current Status And Future](#)
- [Hibernation in Linux 2.6.29](#)
- [System Sleep States — The Linux Kernel documentation](#)
- [Power management/Suspend and hibernate – ArchWiki](#)
- [Hibernation - Debian Wiki](#)
- [Demystifying Hibernation in Linux: Why Isn't It Intuitive?](#)
- [kernel\\_lockdown\(7\) - Linux manual page](#)
- [Hibernation overview - Azure Virtual Machines | Microsoft Learn](#)
- [microsoft/hibernation-setup-tool: Tool to set up a Linux computer to hibernate](#)

THANK YOU!

